

## Model Development Phase Template

Date	18 June 2024
Team ID	SWTID1749713922
Project Title	Early Prediction for Chronic Kidney Disease Detection: A Progressive Approach to Health Management
Maximum Marks	4 Marks

### Initial Model Training Code, Model Validation and Evaluation Report

The model validation and evaluation report will include classification reports, accuracy, and confusion matrices for multiple models, presented through respective screenshots.

### Initial Model Training Code:

```
def run_healthcare_ml_pipeline_hyp(x_train, x_test, y_train, y_test):
    """
    Complete healthcare ML pipeline with enhanced hyperparameter tuning and overfitting detection
    """
    print(f"\n{'='*80}")
    print("HEALTHCARE ML ANALYSIS PIPELINE")
    print("Kidney Disease Detection - Enhanced with Hyperparameter Tuning")
    print(f"{'='*80}")

    # Convert target to binary if needed
    if hasattr(y_train, 'dtype') and y_train.dtype == 'object':
        unique_classes = np.unique(y_train)
        if len(unique_classes) == 2:
            y_train_processed = (y_train == unique_classes[0]).astype(int)
            y_test_processed = (y_test == unique_classes[0]).astype(int)
        else:
            from sklearn.preprocessing import LabelEncoder
            le = LabelEncoder()
            y_train_processed = le.fit_transform(y_train)
            y_test_processed = le.transform(y_test)
    else:
        y_train_processed = y_train
        y_test_processed = y_test
```

```
print(f"Dataset Info:")
print(f"    Training samples: {X_train.shape[0]}")
print(f"    Test samples: {X_test.shape[0]}")
print(f"    Features: {X_train.shape[1]}")
print(f"    Classes: {len(np.unique(y_train_processed))}")
print(f"    Class distribution: {dict(zip(*np.unique(y_train_processed, return_counts=True)))}")

# Create hyperparameter-tuned models
models, tuning_summary = create_tuned_models_dict(X_train, y_train_processed)
print(f"\n Testing {len(models)} different classification algorithms...")
print(f"    {len(tuning_summary)} models have been hyperparameter-tuned")
print(f"    {len(models) - len(tuning_summary)} baseline models included")

# Print hyperparameter tuning summary
if tuning_summary:
    print_hyperparameter_tuning_summary(tuning_summary)

# Run comprehensive analysis with enhanced overfitting detection
results_list, cv_results, overfitting_summary = detect_overfitting_comprehensive_enhanced(
    X_train, X_test, y_train_processed, y_test_processed, models
)

# Convert results list to dictionary for easier access
results_dict = {}
for result in results_list:
    model_name = result['Model']
    results_dict[model_name] = result
```

```
print(f"\n{'='*80}")
print("INDIVIDUAL MODEL ANALYSIS WITH VISUALIZATIONS")
print(f"\n{'='*80}")

# Dictionary to store trained models for plotting
trained_models = {}

# Individual model analysis with plotting
for model_name, model_results in results_dict.items():
    if model_results is None:
        continue
    print(f"\n ANALYZING: {model_name}")
    print("-" * 60)

    try:
        # Get model instance
        if model_name not in models:
            print(f" Model {model_name} not found in models dictionary, skipping...")
            continue

        model = models[model_name]

        # Train the model
        model.fit(X_train, y_train_processed)
        trained_models[model_name] = model

        # Make predictions
        y_pred = model.predict(X_test)
```

```
# Get prediction probabilities if available
if hasattr(model, 'predict_proba'):
    y_pred_proba = model.predict_proba(X_test)
    if len(np.unique(y_test_processed)) == 2:
        y_pred_proba = y_pred_proba[:, 1]
    else:
        y_pred_proba = y_pred_proba.max(axis=1)
elif hasattr(model, 'decision_function'):
    y_pred_proba = model.decision_function(X_test)
    # Normalize decision function scores to [0,1] for binary classification
    if len(np.unique(y_test_processed)) == 2:
        y_pred_proba = (y_pred_proba - y_pred_proba.min()) / (y_pred_proba.max() - y_pred_proba.min())
else:
    y_pred_proba = None

print(f" Model Performance:")
# Access metrics from model_results dictionary
print(f" Test Accuracy: {model_results.get('Test Accuracy', 0):.4f}")
print(f" Precision: {model_results.get('Precision', 0):.4f}")
print(f" Recall: {model_results.get('Recall', 0):.4f}")
print(f" F1 Score: {model_results.get('F1 Score', 0):.4f}")
print(f" ROC AUC: {model_results.get('ROC AUC', 0):.4f}")

# Display hyperparameter tuning results if available
if model_name in tuning_summary:
    tuning_info = tuning_summary[model_name]
    print(f"\n Hyperparameter Tuning Results:")
    print(f" Best Recall Score (CV): {tuning_info['results']['best_score']:.4f}")
    print(f" Overfitting Risk: {tuning_info['results']['overfitting_risk']}")
    print(f" Overfitting Gap: {tuning_info['results']['overfitting_gap']:.4f}")

print(f" Search Method: {tuning_info['results']['search_type']}")

# Show key hyperparameters
key_params = list(tuning_info['params'].items())[:3] # Show first 3 params
if key_params:
    print(f" Key Tuned Parameters:")
    for param, value in key_params:
        print(f" • {param}: {value}")

# Plot confusion matrix
print(f"\n Generating Confusion Matrix...")
plot_confusion_matrix(y_test_processed, y_pred, model_name)

# Plot ROC curve (only for binary classification)
if len(np.unique(y_test_processed)) == 2 and y_pred_proba is not None:
    print(f" Generating ROC Curve...")
    plot_roc_curve(y_test_processed, y_pred_proba, model_name)

    print(f" Generating Precision-Recall Curve...")
    plot_precision_recall_curve(y_test_processed, y_pred_proba, model_name)

# Plot feature importance (if available)
if hasattr(model, 'feature_importances_'):
    print(f" Generating Feature Importance Plot...")
    plot_feature_importance(model, X_train, model_name)
elif hasattr(model, 'coef_') and model.coef_.ndim == 1:
    print(f" Generating Feature Coefficients Plot...")
    # Handle linear model coefficients
```

```
plt.figure(figsize=(10, 6))
if hasattr(X_train, 'columns'):
    coef_series = pd.Series(np.abs(model.coef_), index=X_train.columns)
else:
    coef_series = pd.Series(np.abs(model.coef_), index=[f'Feature_{i}' for i in range(len(model.coef_))])
coef_series.sort_values(ascending=False).head(10).plot(kind='bar')
plt.title(f'Top 10 Feature Coefficients (Absolute) - {model_name}')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.show()

# Validation curve analysis for selected models with hyperparameters
print(f"Generating Validation Curve Analysis...")
if model_name == 'Random Forest':
    validation_curve_analysis_enhanced(
        X_train, y_train_processed, model,
        'n_estimators', [10, 50, 100, 200, 300]
    )
elif model_name == 'XGBoost':
    validation_curve_analysis_enhanced(
        X_train, y_train_processed, model,
        'max_depth', [3, 4, 5, 6, 7, 8]
    )
elif model_name == 'LightGBM':
    validation_curve_analysis_enhanced(
        X_train, y_train_processed, model,
        'num_leaves', [10, 20, 30, 40, 50]
    )
elif 'SVM' in model_name:
    validation_curve_analysis_enhanced(
        X_train, y_train_processed, model,
```

```
        'C', [0.1, 1, 10, 100, 1000]
    )
elif 'Logistic Regression' in model_name:
    validation_curve_analysis_enhanced(
        X_train, y_train_processed, model,
        'C', [0.01, 0.1, 1, 10, 100]
    )

print(f"Completed analysis for {model_name}\n")

except Exception as e:
    print(f"Error analyzing {model_name}: {str(e)}")
    continue

# Print overfitting summary
risk_groups = print_overfitting_summary(overfitting_summary)

# Convert results list to DataFrame for healthcare model selection
results_df = pd.DataFrame(results_list)

# Healthcare-specific model selection
best_model_name, ranked_models = healthcare_model_selection_algorithm(results_df)

# Generate comprehensive model comparison plots
print(f"\n{'='*80}")
print("COMPREHENSIVE MODEL COMPARISON VISUALIZATIONS")
print(f"{'='*80}")
```

```
print(" Generating Top Models Comparison...")
plot_model_comparison(results_df)

# Enhanced hyperparameter tuning summary visualization
if tuning_summary:
    print(f"\n Generating Hyperparameter Tuning Summary Visualization...")
    plot_hyperparameter_tuning_summary(tuning_summary)

# Final recommendations with hyperparameter considerations
print(f"\n{' '*80}")
print(" FINAL HEALTHCARE RECOMMENDATIONS")
print(f"\n{' '*80}")

print(f" RECOMMENDED MODEL: {best_model_name}")
best_stats = ranked_models.iloc[0]
print(f" Healthcare Score: {best_stats['Healthcare_Score']:.4f}")
print(f" Recall (Sensitivity): {best_stats['Recall']:.4f}")
print(f" Precision: {best_stats['Precision']:.4f}")
print(f" F1 Score: {best_stats['F1 Score']:.4f}")
print(f" Overfitting Risk: {best_stats['Overfitting Risk']}")

# Show hyperparameter tuning info for best model if available
if best_model_name in tuning_summary:
    best_tuning = tuning_summary[best_model_name]
    print(f" Hyperparameter Optimization:")
    print(f" • Tuning Method: {best_tuning['results']['search_type']}")
    print(f" • CV Recall Score: {best_tuning['results']['best_score']:.4f}")
    print(f" • Overfitting Gap: {best_tuning['results']['overfitting_gap']:.4f}")
```

```
print(f"\n TOP 3 SAFE MODELS FOR HEALTHCARE:")
safe_models = ranked_models[ranked_models['Overfitting Risk'].isin(['LOW', 'MEDIUM'])].head(3)
for i, (_, row) in enumerate(safe_models.iterrows(), 1):
    risk_indicator = "" if row['Overfitting Risk'] == 'LOW' else ""
    tuning_indicator = "" if row['Model'] in tuning_summary else ""
    print(f" {i}. {tuning_indicator} {row['Model']} (Score: {row['Healthcare_Score']:.4f}, Risk: {risk_indicator}{row['Overfitting Risk']})")

# Models to avoid with hyperparameter tuning context
avoid_models = ranked_models[ranked_models['Overfitting Risk'].isin(['CRITICAL', 'HIGH'])]['Model'].tolist()
if avoid_models:
    print(f"\n MODELS TO AVOID IN HEALTHCARE:")
    for model in avoid_models[:5]: # Show top 5 to avoid
        if model in tuning_summary:
            gap = tuning_summary[model]['results']['overfitting_gap']
            print(f" {model} (Overfitting Gap: {gap:.4f})")
        else:
            print(f" {model}")

# Hyperparameter tuning insights
if tuning_summary:
    print(f"\n HYPERPARAMETER TUNING INSIGHTS:")

    # Count tuned models by risk level
    tuned_risks = {}
    for model_name, info in tuning_summary.items():
        risk = info['results']['overfitting_risk']
        tuned_risks[risk] = tuned_risks.get(risk, 0) + 1
```

```
print(f"    Tuned Models by Risk Level:")
risk_order = ['LOW', 'MEDIUM', 'HIGH', 'CRITICAL']
for risk in risk_order:
    if risk in tuned_risks:
        icon = {'LOW': '●', 'MEDIUM': '●', 'HIGH': '●', 'CRITICAL': '●'}.get(risk, '○')
        print(f"        {icon} {risk}: {tuned_risks[risk]} models")

# Best tuning results
best_tuned_recall = max(tuning_summary.items(), key=lambda x: x[1]['results']['best_score'])
lowest_overfitting_tuned = min(tuning_summary.items(), key=lambda x: x[1]['results']['overfitting_gap'])

print(f"    Best Tuned Recall: {best_tuned_recall[0]} ({best_tuned_recall[1]['results']['best_score']:.4f})")
print(f"    Lowest Overfitting (Tuned): {lowest_overfitting_tuned[0]} (Gap: {lowest_overfitting_tuned[1]['results']['overfitting_gap']:.4f})")

# Additional comprehensive analysis plots
print(f"\n GENERATING ADDITIONAL ANALYSIS PLOTS...")

# Overfitting risk distribution plot
plt.figure(figsize=(15, 10))

# Risk distribution
plt.subplot(2, 3, 1)
risk_counts = ranked_models['Overfitting Risk'].value_counts()
colors = {'LOW': 'green', 'MEDIUM': 'orange', 'HIGH': 'red', 'CRITICAL': 'darkred'}
risk_colors = [colors.get(risk, 'gray') for risk in risk_counts.index]
plt.pie(risk_counts.values, labels=risk_counts.index, autopct='%1.1f%%', colors=risk_colors)
plt.title('Overfitting Risk Distribution')
```

```
# Healthcare scores distribution
plt.subplot(2, 3, 2)
plt.hist(ranked_models['Healthcare_Score'], bins=15, alpha=0.7, color='skyblue', edgecolor='black')
plt.xlabel('Healthcare Score')
plt.ylabel('Number of Models')
plt.title('Healthcare Scores Distribution')
plt.grid(True, alpha=0.3)

# Recall vs Precision scatter plot
plt.subplot(2, 3, 3)
colors_risk = ranked_models['Overfitting Risk'].map(colors)
scatter = plt.scatter(ranked_models['Recall'], ranked_models['Precision'],
                    c=colors_risk, alpha=0.7, s=60, edgecolors='black', linewidth=0.5)
plt.xlabel('Recall (Sensitivity)')
plt.ylabel('Precision')
plt.title('Recall vs Precision (Colored by Risk)')
plt.grid(True, alpha=0.3)

# F1 Score vs CV Stability
plt.subplot(2, 3, 4)
plt.scatter(ranked_models['F1 Score'], ranked_models['CV Std F1'],
            c=colors_risk, alpha=0.7, s=60, edgecolors='black', linewidth=0.5)
plt.xlabel('F1 Score')
plt.ylabel('CV Standard Deviation')
plt.title('Performance vs Stability (Colored by Risk)')
plt.grid(True, alpha=0.3)
```

```
# Hyperparameter tuning comparison (if available)
if tuning_summary:
    plt.subplot(2, 3, 5)
    tuned_models_data = []
    tuned_scores = []
    tuned_gaps = []
    for model_name in ranked_models['Model']:
        if model_name in tuning_summary:
            tuned_models_data.append(model_name[:15]) # Truncate long names
            tuned_scores.append(tuning_summary[model_name]['results']['best_score'])
            tuned_gaps.append(tuning_summary[model_name]['results']['overfitting_gap'])

    if tuned_models_data:
        plt.scatter(tuned_scores, tuned_gaps, alpha=0.7, s=60,
                    c='purple', edgecolors='black', linewidth=0.5)
        plt.xlabel('Tuned CV Recall Score')
        plt.ylabel('Overfitting Gap')
        plt.title('Hyperparameter Tuning Results')
        plt.grid(True, alpha=0.3)

    # Add model names as annotations for top performers
    for i, (score, gap, name) in enumerate(zip(tuned_scores, tuned_gaps, tuned_models_data)):
        if score > np.percentile(tuned_scores, 75) and gap < np.percentile(tuned_gaps, 50):
            plt.annotate(name, (score, gap), xytext=(5, 5),
                        textcoords='offset points', fontsize=8)
```

```
# Model complexity vs performance
plt.subplot(2, 3, 6)
# Create a complexity score based on model type
complexity_map = {
    'Dummy': 1, 'Naive Bayes': 2, 'Logistic Regression': 3, 'LDA': 3, 'QDA': 4,
    'Decision Tree': 4, 'KNN': 4, 'SVM': 5, 'Random Forest': 6, 'Extra Trees': 6,
    'AdaBoost': 6, 'Gradient Boosting': 7, 'XGBoost': 8, 'LightGBM': 8, 'CatBoost': 8,
    'MLP': 9, 'SGD': 3, 'Ridge': 3, 'Bagging': 5
}

complexity_scores = []
for model_name in ranked_models['Model']:
    complexity = 5 # default
    for key, value in complexity_map.items():
        if key.lower() in model_name.lower():
            complexity = value
            break
    complexity_scores.append(complexity)

plt.scatter(complexity_scores, ranked_models['Healthcare_Score'],
            c=colors_risk, alpha=0.7, s=60, edgecolors='black', linewidth=0.5)
plt.xlabel('Model Complexity')
plt.ylabel('Healthcare Score')
plt.title('Complexity vs Healthcare Performance')
plt.grid(True, alpha=0.3)

plt.tight_layout()
plt.show()
```

```
# Summary statistics
print(f"\n SUMMARY STATISTICS:")
print(f"    Total Models Evaluated: {len(results_list)}")
print(f"    Models with Hyperparameter Tuning: {len(tuning_summary)}")
print(f"    Low Risk Models: {len(ranked_models[ranked_models['Overfitting Risk'] == 'LOW'])}")
print(f"    Medium Risk Models: {len(ranked_models[ranked_models['Overfitting Risk'] == 'MEDIUM'])}")
print(f"    High Risk Models: {len(ranked_models[ranked_models['Overfitting Risk'] == 'HIGH'])}")
print(f"    Critical Risk Models: {len(ranked_models[ranked_models['Overfitting Risk'] == 'CRITICAL'])}")
print(f"    Average Healthcare Score: {ranked_models['Healthcare_Score'].mean():.4f}")
print(f"    Average Recall: {ranked_models['Recall'].mean():.4f}")
print(f"    Average Precision: {ranked_models['Precision'].mean():.4f}")

return {
    'results': results_dict,
    'results_list': results_list,
    'cv_results': cv_results,
    'overfitting_summary': overfitting_summary,
    'risk_groups': risk_groups,
    'best_model': best_model_name,
    'ranked_models': ranked_models,
    'safe_models': safe_models,
    'trained_models': trained_models,
    'tuning_summary': tuning_summary,
    'hyperparameter_insights': {
        'tuned_models_count': len(tuning_summary),
        'best_tuned_recall': best_tuned_recall if tuning_summary else None,
        'lowest_overfitting_tuned': lowest_overfitting_tuned if tuning_summary else None,
        'risk_distribution': tuned_risks if tuning_summary else None
    }
}
```



## Model Validation and Evaluation Report:

Model	Classification Report	Accuracy	Confusion Matrix																														
Random Forest	<div>COMPREHENSIVE ANALYSIS: Random Forest</div> <div>=====</div> <div><div><div></div><div>PERFORMANCE METRICS:</div></div><div>Test Accuracy: 0.9250</div><div>Precision: 0.9375</div><div>Recall: 0.9250</div><div>F1 Score: 0.9260</div><div>ROC AUC: 0.9933</div><div>PR AUC: 0.9889</div></div> <div><div><div></div><div>ENHANCED OVERFITTING ANALYSIS:</div></div><div>Train F1 Score: 0.9561</div><div>Test F1 Score: 0.9260</div><div>F1 Performance Gap: 0.0301</div><div>Accuracy Gap: 0.0307</div><div>Overfitting Score: 0.150</div><div>Overfitting Risk: LOW</div></div> <div><div><div></div><div>STRATIFIED K-FOLD CROSS VALIDATION:</div></div><div>CV Scores: ['0.9844', '0.9203', '0.9372', '0.9526', '0.9683']</div><div>CV Mean: 0.9526</div><div>CV Std: 0.0225</div><div>CV Coefficient of Variation: 0.0237</div></div> <div><div><div></div><div>DETAILED CLASSIFICATION REPORT:</div></div><table><thead><tr><th></th><th>precision</th><th>recall</th><th>f1-score</th><th>support</th></tr></thead><tbody><tr><td>0</td><td>1.00</td><td>0.88</td><td>0.94</td><td>50</td></tr><tr><td>1</td><td>0.83</td><td>1.00</td><td>0.91</td><td>30</td></tr><tr><td>accuracy</td><td></td><td></td><td>0.93</td><td>80</td></tr><tr><td>macro avg</td><td>0.92</td><td>0.94</td><td>0.92</td><td>80</td></tr><tr><td>weighted avg</td><td>0.94</td><td>0.93</td><td>0.93</td><td>80</td></tr></tbody></table></div>		precision	recall	f1-score	support	0	1.00	0.88	0.94	50	1	0.83	1.00	0.91	30	accuracy			0.93	80	macro avg	0.92	0.94	0.92	80	weighted avg	0.94	0.93	0.93	80	92.50%	<div>Confusion Matrix - Random Forest</div>
		precision	recall	f1-score	support																												
	0	1.00	0.88	0.94	50																												
	1	0.83	1.00	0.91	30																												
	accuracy			0.93	80																												
macro avg	0.92	0.94	0.92	80																													
weighted avg	0.94	0.93	0.93	80																													
XGBoost	<div>COMPREHENSIVE ANALYSIS: XGBoost</div> <div>=====</div> <div><div><div></div><div>PERFORMANCE METRICS:</div></div><div>Test Accuracy: 0.9625</div><div>Precision: 0.9626</div><div>Recall: 0.9625</div><div>F1 Score: 0.9624</div><div>ROC AUC: 0.9933</div><div>PR AUC: 0.9905</div></div> <div><div><div></div><div>ENHANCED OVERFITTING ANALYSIS:</div></div><div>Train F1 Score: 0.9905</div><div>Test F1 Score: 0.9624</div><div>F1 Performance Gap: 0.0281</div><div>Accuracy Gap: 0.0280</div><div>Overfitting Score: 0.141</div><div>Overfitting Risk: LOW</div></div> <div><div><div></div><div>STRATIFIED K-FOLD CROSS VALIDATION:</div></div><div>CV Scores: ['1.0000', '0.9685', '0.9217', '0.9842', '0.9683']</div><div>CV Mean: 0.9685</div><div>CV Std: 0.0262</div><div>CV Coefficient of Variation: 0.0271</div></div> <div><div><div></div><div>DETAILED CLASSIFICATION REPORT:</div></div><table><thead><tr><th></th><th>precision</th><th>recall</th><th>f1-score</th><th>support</th></tr></thead><tbody><tr><td>0</td><td>0.96</td><td>0.98</td><td>0.97</td><td>50</td></tr><tr><td>1</td><td>0.97</td><td>0.93</td><td>0.95</td><td>30</td></tr><tr><td>accuracy</td><td></td><td></td><td>0.96</td><td>80</td></tr><tr><td>macro avg</td><td>0.96</td><td>0.96</td><td>0.96</td><td>80</td></tr><tr><td>weighted avg</td><td>0.96</td><td>0.96</td><td>0.96</td><td>80</td></tr></tbody></table></div>		precision	recall	f1-score	support	0	0.96	0.98	0.97	50	1	0.97	0.93	0.95	30	accuracy			0.96	80	macro avg	0.96	0.96	0.96	80	weighted avg	0.96	0.96	0.96	80	96.25%	<div>Confusion Matrix - XGBoost</div>
		precision	recall	f1-score	support																												
	0	0.96	0.98	0.97	50																												
	1	0.97	0.93	0.95	30																												
	accuracy			0.96	80																												
macro avg	0.96	0.96	0.96	80																													
weighted avg	0.96	0.96	0.96	80																													

LightG  
BM

#### COMPREHENSIVE ANALYSIS: LightGBM

=====

##### PERFORMANCE METRICS:

Test Accuracy: 1.0000  
Precision: 1.0000  
Recall: 1.0000  
F1 Score: 1.0000  
ROC AUC: 1.0000  
PR AUC: 1.0000

##### ENHANCED OVERFITTING ANALYSIS:

Train F1 Score: 0.9905  
Test F1 Score: 1.0000  
F1 Performance Gap: -0.0095  
Accuracy Gap: -0.0095  
Overfitting Score: 0.800  
Overfitting Risk: CRITICAL  
Overfitting Indicators:  
• Perfect test score detected (99.9%+) - possible data leakage

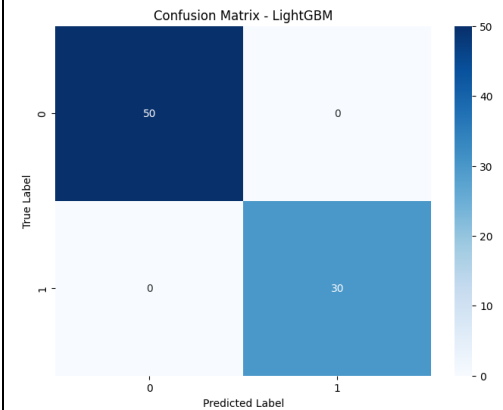
##### STRATIFIED K-FOLD CROSS VALIDATION:

CV Scores: ['0.9844', '0.9842', '0.9842', '0.9842', '0.9841']  
CV Mean: 0.9842  
CV Std: 0.0001  
CV Coefficient of Variation: 0.0001

##### DETAILED CLASSIFICATION REPORT:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	50
1	1.00	1.00	1.00	30
accuracy			1.00	80
macro avg	1.00	1.00	1.00	80
weighted avg	1.00	1.00	1.00	80

100%



Logistic  
Regress  
ion

#### COMPREHENSIVE ANALYSIS: Logistic Regression

=====

##### PERFORMANCE METRICS:

Test Accuracy: 1.0000  
Precision: 1.0000  
Recall: 1.0000  
F1 Score: 1.0000  
ROC AUC: 1.0000  
PR AUC: 1.0000

##### ENHANCED OVERFITTING ANALYSIS:

Train F1 Score: 0.9968  
Test F1 Score: 1.0000  
F1 Performance Gap: -0.0032  
Accuracy Gap: -0.0032  
Overfitting Score: 1.000  
Overfitting Risk: CRITICAL  
Overfitting Indicators:  
• High confidence predictions: 87.0% of training predictions are >95% or <5% confident  
• Perfect test score detected (99.9%+) - possible data leakage


##### STRATIFIED K-FOLD CROSS VALIDATION:

CV Scores: ['0.9844', '1.0000', '0.9842', '0.9842', '0.9842']  
CV Mean: 0.9874  
CV Std: 0.0063  
CV Coefficient of Variation: 0.0064

##### DETAILED CLASSIFICATION REPORT:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	50
1	1.00	1.00	1.00	30
accuracy			1.00	80
macro avg	1.00	1.00	1.00	80
weighted avg	1.00	1.00	1.00	80

100%



## SVM (RBF)

### COMPREHENSIVE ANALYSIS: SVM (RBF)

#### PERFORMANCE METRICS:

Test Accuracy: 1.0000  
Precision: 1.0000  
Recall: 1.0000  
F1 Score: 1.0000  
ROC AUC: 1.0000  
PR AUC: 1.0000

#### ENHANCED OVERFITTING ANALYSIS:

Train F1 Score: 1.0000  
Test F1 Score: 1.0000  
F1 Performance Gap: 0.0000  
Accuracy Gap: 0.0000  
Overfitting Score: 1.000  
Overfitting Risk: CRITICAL  
Overfitting Indicators:  

- High confidence predictions: 97.5% of training predictions are >95% or <5% confident
- Perfect training score detected (99.9%+) - highly suspicious
- Perfect test score detected (99.9%+) - possible data leakage

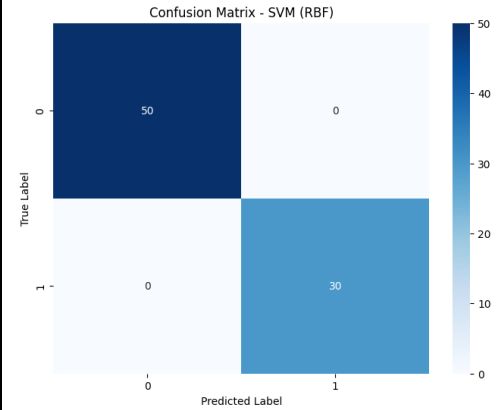
#### STRATIFIED K-FOLD CROSS VALIDATION:

CV Scores: ['0.9844', '0.9842', '1.0000', '1.0000', '1.0000']  
CV Mean: 0.9937  
CV Std: 0.0077  
CV Coefficient of Variation: 0.0077

#### DETAILED CLASSIFICATION REPORT:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	50
1	1.00	1.00	1.00	30
accuracy			1.00	80
macro avg	1.00	1.00	1.00	80
weighted avg	1.00	1.00	1.00	80

100%



## Gradient Boosting

### COMPREHENSIVE ANALYSIS: Gradient Boosting

#### PERFORMANCE METRICS:

Test Accuracy: 0.9500  
Precision: 0.9500  
Recall: 0.9500  
F1 Score: 0.9500  
ROC AUC: 0.9947  
PR AUC: 0.9920

#### ENHANCED OVERFITTING ANALYSIS:

Train F1 Score: 0.9874  
Test F1 Score: 0.9500  
F1 Performance Gap: 0.0374  
Accuracy Gap: 0.0373  
Overfitting Score: 0.187  
Overfitting Risk: LOW

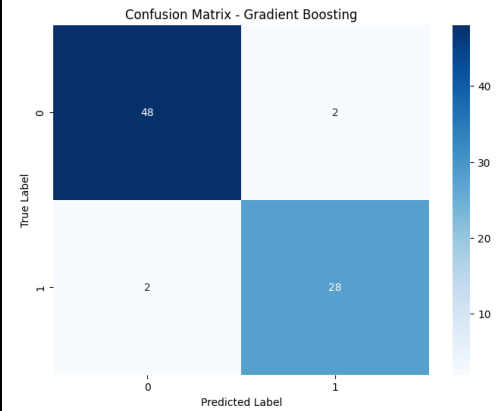
#### STRATIFIED K-FOLD CROSS VALIDATION:

CV Scores: ['0.9844', '0.9842', '0.9528', '0.9842', '0.9841']  
CV Mean: 0.9779  
CV Std: 0.0126  
CV Coefficient of Variation: 0.0128

#### DETAILED CLASSIFICATION REPORT:

	precision	recall	f1-score	support
0	0.96	0.96	0.96	50
1	0.93	0.93	0.93	30
accuracy			0.95	80
macro avg	0.95	0.95	0.95	80
weighted avg	0.95	0.95	0.95	80

95%



CatBoost

#### COMPREHENSIVE ANALYSIS: CatBoost

**PERFORMANCE METRICS:**  
 Test Accuracy: 0.9625  
 Precision: 0.9659  
 Recall: 0.9625  
 F1 Score: 0.9628  
 ROC AUC: 0.9973  
 PR AUC: 0.9958


**ENHANCED OVERFITTING ANALYSIS:**  
 Train F1 Score: 0.9874  
 Test F1 Score: 0.9628  
 F1 Performance Gap: 0.0246  
 Accuracy Gap: 0.0248  
 Overfitting Score: 0.123  
 Overfitting Risk: LOW

**STRATIFIED K-FOLD CROSS VALIDATION:**  
 CV Scores: ['1.0000', '0.9842', '0.9842', '0.9685', '0.9526']  
 CV Mean: 0.9779  
 CV Std: 0.0161  
 CV Coefficient of Variation: 0.0165

**DETAILED CLASSIFICATION REPORT:**

	precision	recall	f1-score	support
0	1.00	0.94	0.97	50
1	0.91	1.00	0.95	30
accuracy			0.96	80
macro avg	0.95	0.97	0.96	80
weighted avg	0.97	0.96	0.96	80

96.25%



Decision Tree

#### COMPREHENSIVE ANALYSIS: Decision Tree

**PERFORMANCE METRICS:**  
 Test Accuracy: 0.8750  
 Precision: 0.8883  
 Recall: 0.8750  
 F1 Score: 0.8767  
 ROC AUC: 0.8867  
 PR AUC: 0.7509


**ENHANCED OVERFITTING ANALYSIS:**  
 Train F1 Score: 0.9561  
 Test F1 Score: 0.8767  
 F1 Performance Gap: 0.0794  
 Accuracy Gap: 0.0807  
 Overfitting Score: 0.397  
 Overfitting Risk: MEDIUM  
 Overfitting Indicators:  
 • Large F1 gap: 0.079 (train: 0.956, test: 0.877)  
 • Large accuracy gap: 0.081

**STRATIFIED K-FOLD CROSS VALIDATION:**  
 CV Scores: ['0.9690', '0.9685', '0.9217', '0.9685', '0.9369']  
 CV Mean: 0.9529  
 CV Std: 0.0199  
 CV Coefficient of Variation: 0.0209

**DETAILED CLASSIFICATION REPORT:**

	precision	recall	f1-score	support
0	0.95	0.84	0.89	50
1	0.78	0.93	0.85	30
accuracy			0.88	80
macro avg	0.87	0.89	0.87	80
weighted avg	0.89	0.88	0.88	80

87.50%



K-Nearest Neighbors

#### COMPREHENSIVE ANALYSIS: K-Nearest Neighbors

**PERFORMANCE METRICS:**  
 Test Accuracy: 1.0000  
 Precision: 1.0000  
 Recall: 1.0000  
 F1 Score: 1.0000  
 ROC AUC: 1.0000  
 PR AUC: 1.0000

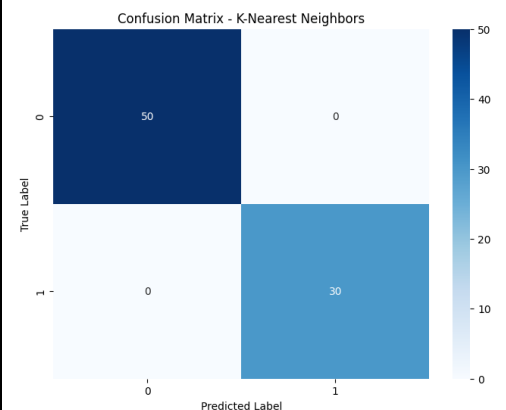
**ENHANCED OVERFITTING ANALYSIS:**  
 Train F1 Score: 1.0000  
 Test F1 Score: 1.0000  
 F1 Performance Gap: 0.0000  
 Accuracy Gap: 0.0000  
 Overfitting Score: 1.000  
 Overfitting Risk: CRITICAL  
 Overfitting Indicators:  
 • High confidence predictions: 100.0% of training predictions are >95% or <5% confident  
 • Perfect training score detected (99.9%+) - highly suspicious  
 • Perfect test score detected (99.9%+) - possible data leakage

**STRATIFIED K-FOLD CROSS VALIDATION:**  
 CV Scores: ['0.9844', '0.9685', '0.9528', '0.9842', '1.0000']  
 CV Mean: 0.9780  
 CV Std: 0.0161  
 CV Coefficient of Variation: 0.0164

**DETAILED CLASSIFICATION REPORT:**

	precision	recall	f1-score	support
0	1.00	1.00	1.00	50
1	1.00	1.00	1.00	30
accuracy			1.00	80
macro avg	1.00	1.00	1.00	80
weighted avg	1.00	1.00	1.00	80

100%



## Linear Discriminant Analysis

### COMPREHENSIVE ANALYSIS: Linear Discriminant Analysis

#### PERFORMANCE METRICS:

Test Accuracy: 1.0000  
Precision: 1.0000  
Recall: 1.0000  
F1 Score: 1.0000  
ROC AUC: 1.0000  
PR AUC: 1.0000

#### ENHANCED OVERFITTING ANALYSIS:

Train F1 Score: 0.9779  
Test F1 Score: 1.0000  
F1 Performance Gap: -0.0221  
Accuracy Gap: -0.0222  
Overfitting Score: 1.000  
Overfitting Risk: CRITICAL  
Overfitting Indicators:  
• High confidence predictions: 93.0% of training predictions are >95% or <5% confident  
• Perfect test score detected (99.9%+) - possible data leakage


#### STRATIFIED K-FOLD CROSS VALIDATION:

CV Scores: ['1.0000', '0.9842', '0.9528', '0.9369', '1.0000']  
CV Mean: 0.9748  
CV Std: 0.0256  
CV Coefficient of Variation: 0.0263

#### DETAILED CLASSIFICATION REPORT:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	50
1	1.00	1.00	1.00	30
accuracy			1.00	80
macro avg	1.00	1.00	1.00	80
weighted avg	1.00	1.00	1.00	80

100%



## Multi-Layer Perceptron

### COMPREHENSIVE ANALYSIS: Multi-layer Perceptron

#### PERFORMANCE METRICS:

Test Accuracy: 0.9875  
Precision: 0.9879  
Recall: 0.9875  
F1 Score: 0.9875  
ROC AUC: 1.0000  
PR AUC: 1.0000

#### ENHANCED OVERFITTING ANALYSIS:

Train F1 Score: 0.9717  
Test F1 Score: 0.9875  
F1 Performance Gap: -0.0158  
Accuracy Gap: -0.0160  
Overfitting Score: 0.000  
Overfitting Risk: LOW

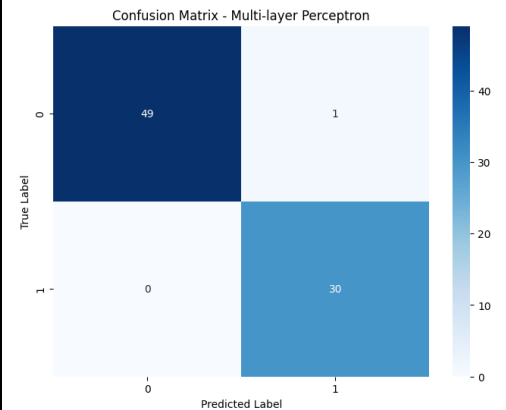
#### STRATIFIED K-FOLD CROSS VALIDATION:

CV Scores: ['0.9690', '0.9528', '0.9685', '0.9685', '0.9842']  
CV Mean: 0.9686  
CV Std: 0.0099  
CV Coefficient of Variation: 0.0102

#### DETAILED CLASSIFICATION REPORT:

	precision	recall	f1-score	support
0	1.00	0.98	0.99	50
1	0.97	1.00	0.98	30
accuracy			0.99	80
macro avg	0.98	0.99	0.99	80
weighted avg	0.99	0.99	0.99	80

98.75%



AdaBoost

#### COMPREHENSIVE ANALYSIS: AdaBoost

##### PERFORMANCE METRICS:

Test Accuracy: 0.8750  
Precision: 0.8883  
Recall: 0.8750  
F1 Score: 0.8767  
ROC AUC: 0.8867  
PR AUC: 0.7509

##### ENHANCED OVERFITTING ANALYSIS:

Train F1 Score: 0.9561  
Test F1 Score: 0.8767  
F1 Performance Gap: 0.0794  
Accuracy Gap: 0.0807  
Overfitting Score: 0.397  
Overfitting Risk: MEDIUM  
Overfitting Indicators:  
• Large F1 gap: 0.079 (train: 0.956, test: 0.877)  
• Large accuracy gap: 0.081

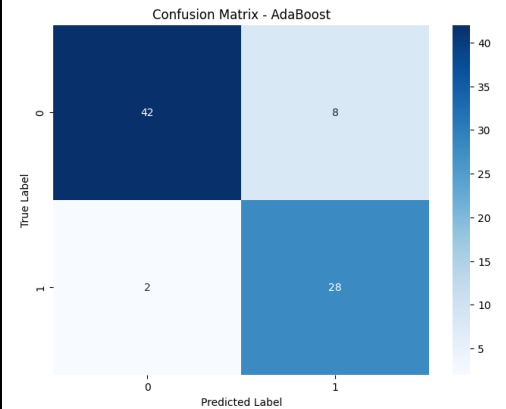
##### STRATIFIED K-FOLD CROSS VALIDATION:

CV Scores: ['0.9690', '0.9685', '0.9217', '0.9685', '0.9369']  
CV Mean: 0.9529  
CV Std: 0.0199  
CV Coefficient of Variation: 0.0209

##### DETAILED CLASSIFICATION REPORT:

	precision	recall	f1-score	support
0	0.95	0.84	0.89	50
1	0.78	0.93	0.85	30
accuracy			0.88	80
macro avg	0.87	0.89	0.87	80
weighted avg	0.89	0.88	0.88	80

87.50%



SGD Classifier

#### COMPREHENSIVE ANALYSIS: SGD Classifier

##### PERFORMANCE METRICS:

Test Accuracy: 1.0000  
Precision: 1.0000  
Recall: 1.0000  
F1 Score: 1.0000  
ROC AUC: 1.0000  
PR AUC: 1.0000

##### ENHANCED OVERFITTING ANALYSIS:

Train F1 Score: 0.9937  
Test F1 Score: 1.0000  
F1 Performance Gap: -0.0063  
Accuracy Gap: -0.0063  
Overfitting Score: 1.000  
Overfitting Risk: CRITICAL  
Overfitting Indicators:  
• High confidence predictions: 97.2% of training predictions are >95% or <5% confident  
• Perfect test score detected (99.9%+) - possible data leakage

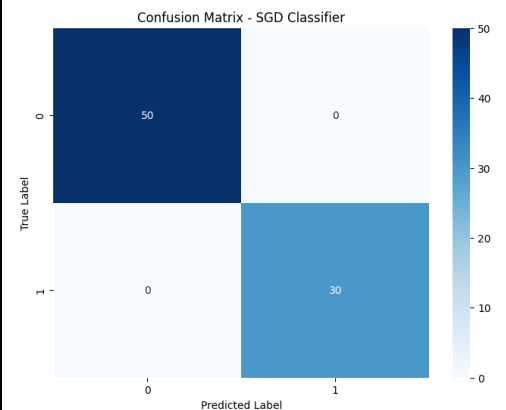
##### STRATIFIED K-FOLD CROSS VALIDATION:

CV Scores: ['0.9844', '1.0000', '0.9842', '0.9683', '0.9842']  
CV Mean: 0.9842  
CV Std: 0.0100  
CV Coefficient of Variation: 0.0102

##### DETAILED CLASSIFICATION REPORT:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	50
1	1.00	1.00	1.00	30
accuracy			1.00	80
macro avg	1.00	1.00	1.00	80
weighted avg	1.00	1.00	1.00	80

100%



## Gaussian Naïve Bayes

### COMPREHENSIVE ANALYSIS: Gaussian Naïve Bayes

#### PERFORMANCE METRICS:

Test Accuracy: 0.9000  
Precision: 0.9211  
Recall: 0.9000  
F1 Score: 0.9015  
ROC AUC: 1.0000  
PR AUC: 1.0000

#### ENHANCED OVERFITTING ANALYSIS:

Train F1 Score: 0.9219  
Test F1 Score: 0.9015  
F1 Performance Gap: 0.0204  
Accuracy Gap: 0.0209  
Overfitting Score: 0.402  
Overfitting Risk: MEDIUM  
Overfitting Indicators:  
• High confidence predictions: 99.7% of training predictions are >95% or <5% confident

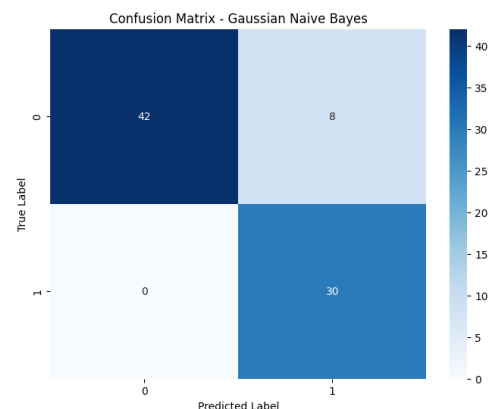
#### STRATIFIED K-FOLD CROSS VALIDATION:

CV Scores: ['0.8616', '0.9372', '0.9528', '0.9061', '0.9528']  
CV Mean: 0.9221  
CV Std: 0.0348  
CV Coefficient of Variation: 0.0377

#### DETAILED CLASSIFICATION REPORT:

	precision	recall	f1-score	support
0	1.00	0.84	0.91	50
1	0.79	1.00	0.88	30
accuracy			0.90	80
macro avg	0.89	0.92	0.90	80
weighted avg	0.92	0.90	0.90	80

90%



## Bernoulli Naïve Bayes

### COMPREHENSIVE ANALYSIS: Bernoulli Naïve Bayes

#### PERFORMANCE METRICS:

Test Accuracy: 0.9750  
Precision: 0.9766  
Recall: 0.9750  
F1 Score: 0.9751  
ROC AUC: 1.0000  
PR AUC: 1.0000

#### ENHANCED OVERFITTING ANALYSIS:

Train F1 Score: 0.9874  
Test F1 Score: 0.9751  
F1 Performance Gap: 0.0122  
Accuracy Gap: 0.0123  
Overfitting Score: 0.361  
Overfitting Risk: MEDIUM  
Overfitting Indicators:  
• High confidence predictions: 97.2% of training predictions are >95% or <5% confident

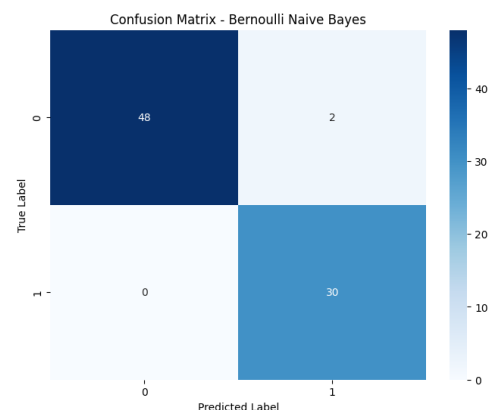
#### STRATIFIED K-FOLD CROSS VALIDATION:

CV Scores: ['1.0000', '0.9842', '0.9842', '0.9685', '0.9841']  
CV Mean: 0.9842  
CV Std: 0.0100  
CV Coefficient of Variation: 0.0101

#### DETAILED CLASSIFICATION REPORT:

	precision	recall	f1-score	support
0	1.00	0.96	0.98	50
1	0.94	1.00	0.97	30
accuracy			0.97	80
macro avg	0.97	0.98	0.97	80
weighted avg	0.98	0.97	0.98	80

97.50%





## Dummy Classifier (Stratified)

### COMPREHENSIVE ANALYSIS: Dummy Classifier (Stratified)

#### PERFORMANCE METRICS:

Test Accuracy: 0.5000  
Precision: 0.4858  
Recall: 0.5000  
F1 Score: 0.4918  
ROC AUC: 0.4533  
PR AUC: 0.3571

#### ENHANCED OVERFITTING ANALYSIS:

Train F1 Score: 0.5709  
Test F1 Score: 0.4918  
F1 Performance Gap: 0.0792  
Accuracy Gap: 0.0696  
Overfitting Score: 1.000  
Overfitting Risk: CRITICAL  
Overfitting Indicators:  
• High confidence predictions: 100.0% of training predictions are >95% or <5% confident  
• Low test performance (0.492) with large gap (0.079) suggests learning noise  
• Large F1 gap: 0.079 (train: 0.571, test: 0.492)  
• Large accuracy gap: 0.070

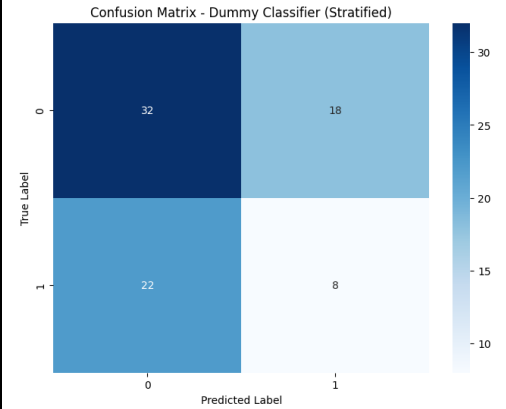
#### STRATIFIED K-FOLD CROSS VALIDATION:

CV Scores: ['0.5982', '0.5594', '0.4941', '0.5594', '0.4615']  
CV Mean: 0.5345  
CV Std: 0.0495  
CV Coefficient of Variation: 0.0926

#### DETAILED CLASSIFICATION REPORT:

	precision	recall	f1-score	support
0	0.59	0.64	0.62	50
1	0.31	0.27	0.29	30
accuracy			0.50	80
macro avg	0.45	0.45	0.45	80
weighted avg	0.49	0.50	0.49	80

50%



## Dummy Classifier (Most Frequent)

### COMPREHENSIVE ANALYSIS: Dummy Classifier (Most Frequent)

#### PERFORMANCE METRICS:

Test Accuracy: 0.6250  
Precision: 0.3906  
Recall: 0.6250  
F1 Score: 0.4808  
ROC AUC: 0.5000  
PR AUC: 0.3750

#### ENHANCED OVERFITTING ANALYSIS:

Train F1 Score: 0.4749  
Test F1 Score: 0.4808  
F1 Performance Gap: -0.0059  
Accuracy Gap: -0.0047  
Overfitting Score: 0.300  
Overfitting Risk: MEDIUM  
Overfitting Indicators:  
• High confidence predictions: 100.0% of training predictions are >95% or <5% confident

#### STRATIFIED K-FOLD CROSS VALIDATION:

CV Scores: ['0.4808', '0.4734', '0.4734', '0.4734', '0.4734']  
CV Mean: 0.4749  
CV Std: 0.0030  
CV Coefficient of Variation: 0.0062

#### DETAILED CLASSIFICATION REPORT:

	precision	recall	f1-score	support
0	0.62	1.00	0.77	50
1	0.00	0.00	0.00	30
accuracy			0.62	80
macro avg	0.31	0.50	0.38	80
weighted avg	0.39	0.62	0.48	80

62.50%

